



Marginal likelihood gap over Gibbs iterations
Computed fresh each step with current η and Θ
EV



non-EV (bad)



non-EV (good)

